

UNIVERSITY NAME

DOCTORAL THESIS

---

# Thesis Title

---

*Author:*

Jan Alexander

FREUDENTHAL

*Supervisor:*

Prof. Arthur KORTE

*A thesis submitted in fulfillment of the requirements*

*for the degree of Ph.D.*

*in the*

Research Group Name

Department or School Name

October 18, 2019

# Declaration of Authorship

I, Jan Alexander FREUDENTHAL, declare that this thesis titled, “Thesis Title” and the work presented in it are my own. I confirm that:

- This work was done wholly or mainly while in candidature for a research degree at this University.
- Where any part of this thesis has previously been submitted for a degree or any other qualification at this University or any other institution, this has been clearly stated.
- Where I have consulted the published work of others, this is always clearly attributed.
- Where I have quoted from the work of others, the source is always given. With the exception of such quotations, this thesis is entirely my own work.
- I have acknowledged all main sources of help.
- Where the thesis is based on work done by myself jointly with others, I have made clear exactly what was done by others and what I have contributed myself.

Signed:

---

Date:

---



*“Thanks to my solid academic training, today I can write hundreds of words on virtually any topic without possessing a shred of information, which is how I got a good job in journalism.”*

Dave Barry



UNIVERSITY NAME

# *Abstract*

Faculty Name

Department or School Name

Ph.D.

**Thesis Title**

by Jan Alexander FREUDENTHAL

The Thesis Abstract is written here (and usually kept to just this page). The page is kept centered vertically so can expand into the blank space above the title too...





# *Acknowledgements*

The acknowledgments and the people to thank go here, don't forget to include your project advisor. . .



# Contents

<b>Declaration of Authorship</b>	<b>iii</b>
<b>Abstract</b>	<b>vii</b>
<b>Acknowledgements</b>	<b>ix</b>
<b>1 Benchmarking of Chloroplast Genome Assembly tools</b>	<b>1</b>
1.1 Introduction . . . . .	1
1.2 Material and Methods . . . . .	1
1.3 Results . . . . .	1
1.3.1 Performance metrics . . . . .	3
1.3.2 Qualitative . . . . .	3
1.3.3 Simulated data . . . . .	3
1.3.4 Real data sets . . . . .	3
1.3.5 Consistency . . . . .	3
1.4 Disucssion . . . . .	3
<b>2 Understanding the hapoltype structure of Arabidopisis thaliana</b>	<b>5</b>
2.1 Introduction . . . . .	5
2.2 Haplotyping of A. thaliana . . . . .	5
2.3 Results . . . . .	5
2.4 Disucssion . . . . .	5
<b>3 GWAS Flow a gpu-accelerated software for large-scale genome-wide association studies</b>	<b>7</b>

3.1	Introduction . . . . .	7
3.2	Methods . . . . .	9
	GWAS Model . . . . .	9
	The GWAS-Flow Software . . . . .	9
	Calculation of permutation-based thresholds for GWAS . . . . .	10
	Benchmarking . . . . .	11
3.3	Results . . . . .	12
3.4	Disucssion . . . . .	13
<b>4</b>	<b>Genomic prediction of phenotypic values of quantitative traits using Artificial neural networks</b>	<b>17</b>
4.1	Introduction . . . . .	17
4.2	Material . . . . .	20
4.3	Methods . . . . .	20
4.4	Results . . . . .	20
4.5	Disucssion . . . . .	20
<b>A</b>	<b>Source code GWAS-Flow</b>	<b>21</b>
A.1	gwas.py . . . . .	21
A.2	main.py . . . . .	24
A.3	herit.py . . . . .	29
	<b>Bibliography</b>	<b>31</b>

# List of Figures

1.1	An Electron . . . . .	2
3.1	Computation time vs number of markers . . . . .	13
3.2	Computations time vs accessions . . . . .	14
3.3	Computational time of GWA Analyses on real <i>A. thaliana</i> data sets . . . . .	16



# List of Tables

1.1	The effects of treatments X and Y on the four groups studied. .	2
-----	---	---





# List of Abbreviations

<b>ANN</b>	<b>Artificial Neural Network</b>
<b>BLUB</b>	<b>Best Linear Unbiased Predictor</b>
<b>BLUE</b>	<b>Best Linear Unbiased Estimator</b>
<b>CPU</b>	<b>Core Processing Unit</b>
<b>EMMA</b>	<b>Efficient Mixed Model Associations</b>
<b>FCL</b>	<b>Fully Connected Layer</b>
<b>GP</b>	<b>Genomic Prediction</b>
<b>GPU</b>	<b>Graphical Processing Unit</b>
<b>GS</b>	<b>Genomic Selection</b>
<b>GWAIS</b>	<b>Genome Wide Interaction Association Studies</b>
<b>GWAS</b>	<b>Genome Wide Association Studies</b>
<b>HDF</b>	<b>Hierarchical Data Format</b>
<b>LCL</b>	<b>Locally Connected Layer</b>
<b>LD</b>	<b>Linkage Disequilibrium</b>
<b>LMM</b>	<b>Linear Mixed Model</b>
<b>MLP</b>	<b>Multi Layer Perceptron</b>
<b>ML</b>	<b>Machine Learning</b>
<b>QTL</b>	<b>Quantitative Trait Locus</b>
<b>RKHS</b>	<b>Reproducing Kernel Hilbert Spaces</b>
<b>RSS</b>	<b>Residual Sum of Squares</b>
<b>SNP</b>	<b>Single Nucleotide Polymorphism</b>
<b>TRN</b>	<b>TRaiNing subset</b>
<b>TST</b>	<b>TeSTing subset</b>

<b>WGS</b>	<b>Whole Genome Sequencing</b>
<b>LSC</b>	<b>Large Single Copy</b>
<b>SSC</b>	<b>Small Single Copy</b>
<b>IR</b>	<b>Inverted Repeat</b>
<b>DNA</b>	<b>DeoxyriboNucleic Acid</b>
<b>DNA</b>	<b>RiboNucleic Acid</b>
<b>GUI</b>	<b>Graphical User Interface</b>

*For/Dedicated to/To my...*



# Chapter 1

## Benchmarking of Chloroplast Genome Assembly tools

### 1.1 Introduction

Here I will but the introduction to from the paper

### 1.2 Material and Methods

### 1.3 Results

$$score = \frac{1}{4} \cdot \left( cov_{ref} + cov_{qry} + \min \left\{ \frac{cov_{qry}}{cov_{ref}}, \frac{cov_{ref}}{cov_{qry}} \right\} + \frac{1}{n_{contigs}} \right) \cdot 100 \quad (1.1)$$



---

FIGURE 1.1: An electron (artist's impression).

TABLE 1.1: The effects of treatments X and Y on the four groups studied.

Groups	Treatment X	Treatment Y
1	0.2	0.8
2	0.17	0.7
3	0.24	0.75
4	0.68	0.3

### **1.3.1 Performance metrics**

### **1.3.2 Qualitative**

### **1.3.3 Simulated data**

### **1.3.4 Real data sets**

### **1.3.5 Consistency**

## **1.4 Disucssion**





## **Chapter 2**

# **Understanding the haplotype structure of *Arabidopsis thaliana***

### **2.1 Introduction**

### **2.2 Haplotyping of *A. thaliana***

### **2.3 Results**

### **2.4 Discussion**



## Chapter 3

# GWAS Flow a gpu-accelerated software for large-scale genome-wide association studies

### 3.1 Introduction

Genome-wide association studies, pioneered in human genetics Hirschhorn and Daly, 2005 in the last decade, have become the predominant method to detect associations between phenotypes and the genetic variations present in a population. Understanding the genetic architecture of traits and mapping the underlying genomic polymorphisms is of paramount importance for successful breeding both in plants and animals, as well as for studying the genetic risk factors of diseases. Over the last decades, the cost for genotyping have been reduced dramatically. Early GWAS consisted of a few hundred individuals which have been phenotyped and genotyped on a couple of hundreds to thousands of genomic markers. Nowadays, marker density for many species easily exceed millions of genomic polymorphisms. Albeit commonly SNPs are used for association studies, standard GWAS models are flexible to handle different genomic features as input. The *Arabidopsis* 1001

genomes project features for example 1135 sequenced *Arabidopsis thaliana* accessions with over 10 million genomic markers that segregate in the population Alonso-Blanco et al., 2016. Other genome projects also yielded large amounts of genomic data for a substantial amount of individuals, as exemplified in the 1000 genomes project for humans Siva, 2008, the 2000 yeast genomes project or the 3000 rice genomes project Li, Wang, and Zeigler, 2014. Thus, there is an increasing demand for GWAS models that can analyze these data in a reasonable time frame. One critical step of GWAS is to determine the threshold at which an association is termed significant. Classically the conservative Bonferroni threshold is used, which accounts for the number of statistical tests that are performed, while many recent studies try to use significance thresholds that are based on the false-discovery rate (FDR) Storey and Tibshirani, 2003. An alternative approach are permutation-based thresholds Che et al., 2014. Permutation-based thresholds estimate the significance by shuffling phenotypes and genotypes before each GWAS run, thus any signal left in the data should not have a genetic cause, but might represent model mis-specifications or uneven phenotypic distributions. Typically this process is repeated hundreds to thousands of times and will lead to a distinct threshold for each phenotype analyzed Togninalli et al., 2017. The computational demand of permutation-based thresholds is immense, as per analysis not one, but at least hundreds of GWAS need to be performed. Here the main limitation is the pure computational demand. Thus, faster GWAS models could easily make the estimation of permutation-based thresholds the default choice.

## 3.2 Methods

### GWAS Model

The GWAS model used for GWAS-Flow is based on a fast approximation of the linear-mixed-model described in Kang et al., 2010; Zhang et al., 2010, which estimates the variance components  $\sigma_g$  and  $\sigma_e$  only once in a null model that includes the genetic relationship matrix, but no distinct genetic markers. These components are thereafter used for the tests of each specific marker. Here, the underlying assumption is, that the ratio of these components stays constant, even if distinct genetic markers are included into the GWAS model. This holds true for nearly all markers and only markers which possess a big effect will alter this ratio slightly, where now  $\sigma_g$  would become smaller compared to the null model. Thus, the p-values calculated by the approximation might be a little higher (less significant) for strongly associated markers.

### The GWAS-Flow Software

The GWAS-Flow software was designed to provide a fast and robust GWAS implementation that can easily handle large data and allows to perform permutations in a reasonable time frame. Traditional GWAS implementations that are implemented using Python Van Rossum and Drake Jr, 1995 or R R Core Team, 2019 cannot always meet these demands. We tried to overcome those limitations by using TensorFlow Abadi et al., 2015, a multi-language machine learning framework published and developed by Google. GWAS calculations are composed of a series of matrix computations that can be highly parallelized, and easily integrated into the architecture provided by TensorFlow. Our implementation allows both, the classical parallelization of code on multiple processors (CPUs) and the use of graphical processing units (GPUs). GWAS-Flow is written using the Python TensorFlow API. Data import is done with *pandas* McKinney, 2010 and/or *HDF5* for Python Collette,

2013. Preprocessing of the data (e.g filtering by minor Allele count (MAC)) is performed with *numpy* Oliphant, 2006. Variance components for residual and genomic effects are estimated with a slightly altered function based on the Python package *limix* Lippert et al., 2014. The GWAS model is based on the following linear mixed model that takes into account the effect of every marker with respect to the kinship:

$$Y = \beta_0 + X_i\beta_i + u + \epsilon, u \sim N(0, \sigma_g K), \epsilon \sim N(0, \sigma_e I) \quad (3.1)$$

From this LMM the residual sum of squares for marker  $i$  are calculated as described in 3.2

$$RSS_i = \sum Y - (X_i\beta_0 + I_i\beta_1) \quad (3.2)$$

The residuals are used to calculate a p-value for each marker according to an overall F-test that compares the model including a distinct genetic effect to a model without this genetic effect:

$$F = \frac{RSS_{env} - R1_{full}}{\frac{R1_{full}}{n-3}} \quad (3.3)$$

Apart from the p-values that derive from the F-distribution, GWAS-Flow also report summary statistics, such as the estimated effect size ( $\beta_i$ ) and its standard error for each marker.

### Calculation of permutation-based thresholds for GWAS

To calculate a permutation-based threshold, we essentially perform  $n$  repetitions ( $n > 100$ ) of the GWAS on the same data with the sole difference that before each GWAS we randomize the phenotypic values. Thus any correlation between the phenotype and the genotype will be broken and indeed for over 90% of these analyses the estimated pseudo-heritability is close to zero.

On the other hand, the phenotypic distribution will stay unaltered by this randomization. Hence, any remaining signal in the GWAS has to be of a non-genetic origin and could be caused by e.g. model mis-specifications. Now we take the lowest p-value (after filtering for the desired minor allele count) for each permutation and take the 5% lowest value as the permutation-based threshold for the GWAS.

### Benchmarking

For benchmarking of GWAS-Flow we used data from the *Arabidopsis* 1001 Genomes Project Alonso-Blanco et al., 2016. The genomic data we used were subsets between 10,000 and 100,000 markers. We chose not to include subsets that exceed 100,000 markers, because there is a linear relationship between the number of markers and the computational time demanded, as all markers are tested independently. We used phenotypic data for flowering time at ten degrees (FT10) for *A. thaliana*, published and downloaded from the AraPheno database Seren et al., 2016. We down- and up-sampled sets to generate phenotypes for sets between 100 and 5000 accessions. For each set of phenotypes and markers we ran 10 permutations to assess the computational time needed. All analyses have been performed with a custom R script that has been used previously Togninalli et al., 2017, GWAS-Flow using either a CPU or a GPU architecture and GEMMA Zhou and Stephens, 2012. GEMMA is a fast and efficient implementation of the mixed model that is broadly used to perform GWAS. All calculations were run on the same machine using 16 i9 virtual CPUs. The GPU version ran on an NVIDIA Tesla P100 graphic card. Additionally to the analyses of the simulated data, we compared the times required by GEMMA and both GWAS-Flow implementations for > 200 different real datasets from *A. thaliana* that have been downloaded from the AraPheno Seren et al., 2016 database and have been analyzed with the available fully

imputed genomic dataset of ca. 10 million markers, filtered for a minor allele count greater five.

### 3.3 Results

The two main factors influencing the computational time for GWAS are the number of markers incorporated in such an analysis and the number of different accessions, while the latter has an approximate quadratic effect in classical GWAS implementations Zhou and Stephens, 2012. Figure 1A shows the time demand as a function of the number of accessions used in the analysis with 10,000 markers. The quadratic increase in time demand is clearly visible for the custom R implementation, as well as for the CPU-based GWAS-Flow implementation and *GEMMA*. The GWAS-Flow implementation and *GEMMA* clearly outperforms the R implementation in general, while for a small number of accessions GWAS-Flow is slightly faster than *GEMMA*. For the GPU-based implementation the increase in run-time with larger sample sizes is much less pronounced. While for small ( $< 1,000$  individuals) data, there is no benefit compared to running GWAS-Flow on CPUs or running *GEMMA*, the GPU-version clearly outperforms the other implementations if the number of accessions increases. Figure 1B shows the computational time in relation to the number of markers and a fixed amount of 2000 accessions for the two different GWAS-Flow implementations. Here, a linear relationship is visible in both cases. To show the performance of GWAS-Flow not only for simulated data, we also run both implementations on more than 200 different real datasets downloaded from the AraPheno database. Figure 1C shows the computational time demands for all analyses comparing both GWAS-Flow implementation to *GEMMA*. Here, the CPU-based GWAS-Flow performs comparable to *GEMMA*, while the GPU-based implementation outperforms both,



if the number of accessions is above 500. Importantly all obtained GWAS results (p-values, beta estimates and standard errors of the beta estimates) are nearly (apart from some mathematical inaccuracies) identical between the three different implementations.




FIGURE 3.1: Computational time as a function of the number of genetic markers with constantly 2000 accessions for both GWAS-Flow versions

## 3.4 Disucssion

We made use of recent developments of computational architecture and software to cope with the increasing computational demand in analyzing large




FIGURE 3.2: Computational time as a function of the number of accessions with 10000 markers each.

GWAS datasets. With GWAS-Flow we implemented both, a CPU- and a GPU-based version of the classical linear mixed model commonly used for GWAS. Both implementations outperform custom R scripts on simulated and real data. While the CPU-based version performs nearly identical compared to *GEMMA*, a commonly used GWAS implementation, the GPU-based implementation outperforms both, if the number of individuals, which have been phenotyped, increases. For analyzing big data, here the main limitation would be the RAM of the GPU, but as the individual test for each marker are independent, this can be easily overcome programmatically. The presented GWAS-Flow implementations are markedly faster compared to custom GWAS

scripts and even outperform efficient fast implementations like *GEMMA* in terms of speed. This readily enables the use of permutation-based thresholds, as with *GWAS-Flow* hundred permutations can be performed in a reasonable time even for big data. Thus, it is possible for each analyzed phenotype to create a specific, permutation-based threshold that might present a more realistic scenario. Importantly the permutation-based threshold can be easily adjusted to different minor allele counts, generating different significance thresholds depending on the allele count. This could help to distinguish false and true associations even for rare alleles. *GWAS-Flow* is a versatile and fast software package. Currently *GWAS-Flow* is and will remain under active development to make the software more versatile. This will e.g. include the compatibility with TensorFlow v2.0.0 and enable data input formats, such as PLINK Purcell et al., 2007. The whole framework is flexible, so it is easy to include predefined co-factors e.g. to enable multi-locus models Segura et al., 2012 or account for multi-variate models like the multi-trait mixed model Korte et al., 2012. Standard GWAS are good in detecting additive effects with comparably large effect sizes, but lack the ability to detect epistatic interactions and their influence on complex traits Mckinney and Pajewski, 2012; Korte and Farlow, 2013a. To catch the effects of these gene-by-gene or SNP-by-SNP interactions, a variety of genome-wide association interaction studies (GWAIS) have been developed, thoroughly reviewed in Ritchie and Van Steen, 2018. Here, *GWAS-Flow* might provide a tool that enables to test the full pairwise interaction matrix of all SNPs. Although this might be a statistic nightmare, it now would be computationally feasible.




FIGURE 3.3: Comparison of the computational time for the analyses of  $> 200$  phenotypes from *Arabidopsis thaliana* as a function of the number of accessions for GEMMA and the CPU- and GPU-based version of GWAS-Flow. GWAS was performed with a fully imputed genotype matrix containing 10.7 M markers and a minor allele filter of  $MAC > 5$

## Chapter 4

# Genomic prediction of phenotypic values of quantitative traits using Artificial neural networks

### 4.1 Introduction

Genomic selection (GS) has been successfully applied in animal Gianola and Rosa, 2015, Hayes and Goddard, 2010 and plant breeding Crossa et al., 2010, Desta and Ortiz, 2014, Heffner et al., 2010, Crossa et al., 2017a as well as in medical applications, since it was first reported Hayes and Goddard, 2001a. Since then the repertoire of methods for predicting phenotypic values has increased rapidly e.g. De Los Campos et al., 2009, Habier et al., 2011, Gianola, 2013, Crossa et al., 2017b. The most commonly applied methods include GBLUP and a set of related algorithms known as the bayesian alphabet Gianola et al., 2009. Genomic prediction in general has repeatedly been shown to outperform pedigree-based methods Crossa et al., 2010, Albrecht et al., 2011 and is nowadays used in many plant and animal breeding schemes. It has also been shown that using whole-genome information is superior to using only feature-selected markers with known QTLs for a given trait Bernardo and Yu, 2007, Heffner, Jannink, and Sorrells, 2011 in some cases. A

more recent study Azodi et al., 2019 compared 11 different genomic prediction algorithms with a variety of data sets and found contradicting results, indicating that feature selection can be useful in some cases when the whole genome regression is performed by neural nets. While every new method is a valuable addition to the tool-kits for genomic selection, some fundamental problems remain unsolved, of which the  $n \gg p$  problematic stands out. Usually in genomic selection settings the size of the training population (TRN) with  $n$  phenotypes is substantially smaller than the number of markers ( $p$ ) Fan, Han, and Liu, 2014. Making the number of features immensely large, even when SNP-SNP interactions are not considered. Furthermore each marker is treated as an independent observation neglecting collinearity and linkage disequilibrium (LD). Further difficulties arise through non-additive, epistatic and dominance marker effects. The main problem with epistasis in quantitative genetics is the almost infinite amount of different marker combinations, that cannot be represented within the size of TRN in the thousands, the same problems arise for example in GWA studies Korte and Farlow, 2013b. With already large  $p$  the number of possible additive SNP-SNP interactions potentiates to  $p^{(p-1)}$ . Methods that attempt to overcome those issues are EGBLUP, using an enhanced epistatic kinship matrix and reproducing kernel Hilbert space regression (RKHS) Jiang and Reif, 2015, Martini et al., 2017.

In the past 10 years, due to increasing availability of high performance computational hardware with decreasing costs and parallel development of free easy-to-use software, most prominent being googles library TensorFlow Abadi et al., 2016 and Keras Chollet, 2015, machine learning (ML) has experienced a renaissance. ML is a set of methods and algorithms used widely for regression and classification problems. popular among those are e.g. support vector machines, multi-layer perceptrons (MLP) and convolutional neural networks. The machine learning mimics the architecture of neural networks and are therefore commonly referred to as artificial neural networks (ANN).

Those algorithms have widely been applied in many biological fields Min, Lee, and Yoon, 2017 , Lan et al., 2018, Mamoshina et al., 2016, Angermueller et al., 2016 , Webb, 2018, Rampasek and Goldenberg, 2016.

A variety of studies assessed the usability of ML in genomic prediction González-Camacho et al., 2018, González-Camacho et al., 2016, Ogutu, Piepho, and Schulz-Streeck, 2011, Montesinos-López et al., 2019a, Grinberg, Orhobor, and King, 2018 , Cuevas et al., 2019, Montesinos-López et al., 2019b, Ma et al., 2017, Qiu et al., 2016, González-Camacho et al., 2012 Li et al., 2018. Through all those studies the common denominator is that there is no such thing as a gold standard for genomic prediction. No single algorithm was able to outperform all the others tested in a single of those studies, let alone in all. While the generally aptitude of ML for genomic selection has been repeatedly shown, how no evidence exists that neural networks can outperform or in many cases perform on that same level as mixed-model approaches as GBLUP Hayes and Goddard, 2001b. While in other fields like image classification neural networks have up to 100s of hidden layers He et al., 2016 the commonly used fully-connected networks in genomic prediction of 1 - 3 hidden layers. With 1 layer networks often being the most successful among those. Contradicting to the idea behind machine learning in genomic selection 1 hidden layer networks will be inapt to capture interactions between loci and thus only account for additive effects. As shown in Azodi et al., 2019 convolutional networks perform worse than fully-connected networks in genomic selection, which again is contradicting to other fields where convolutional layers are applied successfully, e.g natural language processing Dos Santos and Gatti, 2014 or medical image analysis Litjens et al., 2017. Instead of using convolutional layers and fully-connected layers only, as show in Pook et al 2019, we also propose to use locally-connected layer in combination with fully-connected layers. While CL and LCL are closely related they have a significant difference. While in CL weights are shared between

neurons in LCLs each neuron as its own weight. This leads to a reduced number of parameters to be trained in the following FCLs, and should therefore theoretically lead to a decrease in overfitting a common problem in machine learning. To evaluate the results of Pook et al. 2019 accomplished with simulated data we used the data sets generated in the scope of the 1001 genome project of *Arabidopsis thaliana* Alonso-Blanco et al., 2016

## 4.2 Material

## 4.3 Methods

## 4.4 Results

## 4.5 Disucssion



## Appendix A

# Source code GWAS-Flow

### A.1 gwas.py

```
1 import os
2 import sys
3 import time
4 import numpy as np
5 import pandas as pd
6 import main
7 import h5py
8
9 # set defaults
10 mac_min = 1
11 batch_size = 500000
12 out_file = "results.csv"
13 m = 'phenotype_value'
14 perm = 1
15 mac_min= 6
16
17 X_file = 'gwas_sample_data/AT_geno.hdf5'
18 Y_file = 'gwas_sample_data/phenotype.csv'
19 K_file = 'gwas_sample_data/kinship_ibs_binary_mac5.h5py'
20
21
22
23 for i in range (1,len(sys.argv),2):
24     if sys.argv[i] == "-x" or sys.argv[i] == "--genotype":
25         X_file = sys.argv[i+1]
```

```

26     elif sys.argv[i] == "-y" or sys.argv[i] == "--phenotype":
27         Y_file = sys.argv[i+1]
28     elif sys.argv[i] == "-k" or sys.argv[i] == "--kinship":
29         K_file = sys.argv[i+1]
30     elif sys.argv[i] == "-m":
31         m = sys.argv[i+1]
32     elif sys.argv[i] == "-a" or sys.argv[i] == "--mac_min":
33         mac_min = int(sys.argv[i+1])
34     elif sys.argv[i] == "-bs" or sys.argv[i] == "--batch-size":
35         batch_size = int(sys.argv[i+1])
36     elif sys.argv[i] == "-p" or sys.argv[i] == "--perm":
37         perm = int(sys.argv[i+1])
38     elif sys.argv[i] == "-o" or sys.argv[i] == "--out":
39         out_file = sys.argv[i+1]
40     elif sys.argv[i] == "-h" or sys.argv[i] == "--help":
41         print("-x , --genotype :file containing marker
information in csv or hdf5 format of size")
42         print("-y , --phenotype: file container phenotype
information in csv format" )
43         print("-k , --kinship : file containing kinship matrix
of size k X k in csv or hdf5 format")
44         print("-m : name of columnn containing the phenotype :
default m = phenotype_value")
45         print("-a , --mac_min : integer specifying the minimum
minor allele count necessary for a marker to be included.
Default a = 1" )
46         print("-bs, --batch-size : integer specifying the number
of markers processed at once. Default -bs 500000" )
47         print("-p , --perm : single integer specifying the
number of permutations. Default 1 == no perm ")
48         print("-o , --out : name of output file. Default -o
results.csv ")
49         print("-h , --help : prints help and command line
options")
50         quit()
51     else:
52         print('unknown option ' + str(sys.argv[i]))
53         quit()

```

```

54
55
56
57 print("parsed commandline args")
58
59 start = time.time()
60
61 X,K,Y_,markers = main.load_and_prepare_data(X_file,Y_file,K_file
        ,m)
62
63
64 ## MAF filterin
65 markers_used , X , macs = main.mac_filter(mac_min,X,markers)
66
67 ## prepare
68 print("Begin performing GWAS on ", Y_file)
69
70 if perm == 1:
71     output = main.gwas(X,K,Y_,batch_size)
72     if( X_file.split(".")[ -1] == 'csv'):
73         chr_pos = np.array(list(map(lambda x : x.split("- "),
markers_used)))
74     else:
75         chr_reg = h5py.File(X_file,'r')['positions'].attrs['
chr_regions']
76         mk_index= np.array(range(len(markers)),dtype=int)[macs
>= mac_min]
77         chr_pos = np.array([list(map(lambda x: sum(x > chr_reg
[: ,1]) + 1, mk_index)), markers_used]).T
78         my_time = np.repeat((time.time()-start),len(chr_pos))
79         pd.DataFrame({
80             'chr' : chr_pos[:,0] ,
81             'pos' : chr_pos[:,1] ,
82             'pval': output[:,0] ,
83             'mac' : np.array(macs[macs >= mac_min],dtype=np.int) ,
84             'eff_size': output[:,1] ,
85             'SE' : output[:,2]}) .to_csv(out_file,index=False)
86 elif perm > 1:

```

```

87     min_pval = []
88     perm_seeds = []
89     my_time = []
90     for i in range(perm):
91         start_perm = time.time()
92         print("Running permutation ", i+1, " of ", perm)
93         my_seed = np.asscalar(np.random.randint(9999, size=1))
94         perm_seeds.append(my_seed)
95         np.random.seed(my_seed)
96         Y_perm = np.random.permutation(Y_)
97         output = main.gwas(X, K, Y_perm, batch_size)
98         min_pval.append(np.min(output[:, 0]))
99         print("Elapsed time for permuatation", i+1, " with p_min"
, min_pval[i], " is", ":", round(time.time() - start_perm, 2))
100         my_time.append(time.time() - start_perm)
101     pd.DataFrame({
102         'time': my_time,
103         'seed': perm_seeds,
104         'min_p': min_pval }).to_csv(out_file, index=False)
105
106 print("done")
107
108 end = time.time()
109 eltime = np.round(end - start, 2)
110
111 if eltime <= 59:
112     print("Total time elapsed", eltime, "seconds")
113 elif eltime > 59 and eltime <= 3600:
114     print("Total time elapsed", np.round(eltime / 60, 2), "
minutes")
115 elif eltime > 3600 :
116     print("Total time elapsed", np.round(eltime / 60 / 60, 2), "
hours")
117
118

```

## A.2 main.py

```

1     import pandas as pd
2     import numpy as np
3     from scipy.stats import f
4     import tensorflow as tf
5     import limix
6     import herit
7     import h5py
8     import limix
9     import multiprocessing as mlt
10
11 def load_and_prepare_data(X_file,Y_file,K_file,m):
12     type_K = K_file.split(".")[1]
13     type_X = X_file.split(".")[1]
14
15     ## load and preprocess genotype matrix
16     Y = pd.read_csv(Y_file,engine='python').sort_values(['
accession_id']).groupby('accession_id').mean()
17     Y = pd.DataFrame({'accession_id' : Y.index, '
phenotype_value' : Y[m]})
18     if type_X == 'hdf5' or type_X == 'h5py' :
19         SNP = h5py.File(X_file,'r')
20         markers= np.asarray(SNP['positions'])
21         acc_X = np.asarray(SNP['accessions'][:],dtype=np.int)
22     elif type_X == 'csv' :
23         X = pd.read_csv(X_file,index_col=0)
24         markers = X.columns.values
25         acc_X = X.index
26         X = np.asarray(X,dtype=np.float32)/2
27     else :
28         sys.exit("Only hdf5, h5py and csv files are supported")
29
30     if type_K == 'hdf5' or type_K == 'h5py':
31         k = h5py.File(K_file,'r')
32         acc_K = np.asarray(k['accessions'][:],dtype=np.int)
33     elif type_K == 'csv':
34         k = pd.read_csv(K_file,index_col=0)
35         acc_K = k.index
36         k = np.array(k, dtype=np.float32)

```

```

37
38     acc_Y = np.asarray(Y[['accession_id']]).flatten()
39     acc_isec = [isec for isec in acc_X if isec in acc_Y]
40
41     idx_acc = list(map(lambda x: x in acc_isec, acc_X))
42     idy_acc = list(map(lambda x: x in acc_isec, acc_Y))
43     idk_acc = list(map(lambda x: x in acc_isec, acc_K))
44
45     Y_ = np.asarray(Y.drop('accession_id',1),dtype=np.float32)[
idy_acc,:]
46
47     if type_X == 'hdf5' or type_X == 'h5py' :
48         X = np.asarray(SNP['snps'][0:(len(SNP['snps'])+1)],,
dtype=np.float32)[: ,idx_acc].T
49         X = X[np.argsort(acc_X[idx_acc]),:]
50         k1 = np.asarray(k['kinship'][:])[idk_acc,:]
51         K = k1[:,idk_acc]
52         K = K[np.argsort(acc_X[idx_acc]),:]
53         K = K[:,np.argsort(acc_X[idx_acc])]
54     else:
55         X = X[idx_acc,:]
56         k1 = k[idk_acc,:]
57         K = k1[:,idk_acc]
58
59
60     print("data has been imported")
61     return X,K,Y_,markers
62
63
64 def mac_filter(mac_min, X, markers):
65     ac1 = np.sum(X,axis=0)
66     ac0 = X.shape[0] - ac1
67     macs = np.minimum(ac1,ac0)
68     markers_used = markers[macs >= mac_min]
69     X = X[:,macs >= mac_min]
70     return markers_used, X, macs
71
72 def gwas(X,K,Y,batch_size):

```

```

73     n_marker = X.shape[1]
74     n = len(Y)
75     ## REML
76     K_stand = (n-1)/np.sum((np.identity(n) - np.ones((n,n))/n) *
77                             K) * K
78     vg, delta, ve = herit.estimate(Y,"normal",K_stand,verbose =
79                                   False)
80     print(" Pseudo-heritability is " , vg / (ve + vg + delta))
81     print(" Performing GWAS on ", n , " phenotypes and ",
82           n_marker , "markers")
83     ## Transform kinship-matrix, phenotypes and estimate
84     intercept
85     Xo = np.ones(K.shape[0]).flatten()
86     M = np.transpose(np.linalg.inv(np.linalg.cholesky(vg *
87                                                       K_stand + ve * np.identity(n))))
88       .astype(np.float32)
89     Y_t = np.sum(np.multiply(np.transpose(M),Y),axis=1).astype(
90               np.float32)
91     int_t = np.sum(np.multiply(np.transpose(M),np.ones(n)),axis
92                     =1).astype(np.float32)
93     ## EMMAX Scan
94     RSS_env = (np.linalg.lstsq(np.reshape(int_t,(n,-1)) , np.
95                               reshape(Y_t,(n,-1)))[1]).astype(np.float32)
96     ## calculate betas and se of betas
97     def stderr(a,M,Y_t2d,int_t):
98         x = tf.stack((int_t,tf.squeeze(tf.matmul(M.T,tf.reshape
99             (a,(n,-1))))),axis=1)
100         coeff = tf.matmul(tf.matmul(tf.linalg.inv(tf.matmul(tf.
101             transpose(x),x)),tf.transpose(x)),Y_t2d)
102         SSE = tf.reduce_sum(tf.math.square(tf.math.subtract(Y_t
103             ,tf.math.add(tf.math.multiply(x[:,1],coeff[0,0]),tf.math.
104             multiply(x[:,1],coeff[1,0])))))
105         SE = tf.math.sqrt(SSE/(471-(1+2)))
106         StdERR = tf.sqrt(tf.linalg.diag_part(tf.math.multiply(
107             SE , tf.linalg.inv(tf.matmul(tf.transpose(x),x))))[1]
108         return tf.stack((coeff[1,0],StdERR))
109     ## calculate residual sum squares
110     def rss(a,M,y,int_t):
111         x_t = tf.reduce_sum(tf.math.multiply(M.T,a),axis=1)

```

```

98         lm_res = tf.linalg.lstsq(tf.transpose(tf.stack((int_t,
x_t),axis=0)),Y_t2d)
99         lm_x = tf.concat((tf.squeeze(lm_res),x_t),axis=0)
100         return tf.reduce_sum(tf.math.square(tf.math.subtract(tf
.squeeze(Y_t2d),tf.math.add(tf.math.multiply(lm_x[1],lm_x
[2:]), tf.multiply(lm_x[0],int_t)))))
101     ## loop over the batches
102     for i in range(int(np.ceil(n_marker/batch_size))):
103         tf.reset_default_graph()
104         if n_marker < batch_size:
105             X_sub = X
106         else:
107             lower_limit = batch_size * i
108             upper_limit = batch_size * i + batch_size
109             if upper_limit <= n_marker :
110                 X_sub = X[:,lower_limit:upper_limit]
111                 print("Working on markers ", lower_limit , " to
", upper_limit, " of ", n_marker )
112             else:
113                 X_sub = X[:,lower_limit:]
114                 print("Working on markers ", lower_limit , " to
", n_marker, " of ", n_marker )
115             config = tf.ConfigProto()
116             n_cores = mlt.cpu_count()
117             config.intra_op_parallelism_threads = n_cores
118             config.inter_op_parallelism_threads = n_cores
119             sess = tf.Session(config=config)
120             Y_t2d = tf.cast(tf.reshape(Y_t,(n,-1)),dtype=tf.float32)
121             y_tensor = tf.convert_to_tensor(Y_t,dtype = tf.float32)
122             StdERR = tf.map_fn(lambda a : stderr(a,M,Y_t2d,int_t),
X_sub.T)
123             R1_full = tf.map_fn(lambda a: rss(a,M,Y_t2d,int_t),
X_sub.T)
124             F_1 = tf.divide(tf.subtract(RSS_env, R1_full),tf.divide(
R1_full,(n-3)))
125             if i == 0 :
126                 output = sess.run(tf.concat([tf.reshape(F_1,(X_sub.
shape[1],-1)),StdERR],axis=1))

```



```

127         else :
128             tmp = sess.run(tf.concat([tf.reshape(F_1,(X_sub.
shape[1],-1)),StdERR],axis=1))
129             output = np.append(output,tmp,axis=0)
130             sess.close()
131             F_dist = output[:,0]
132             pval = 1 - f.cdf(F_dist,1,n-3)
133             output[:,0] = pval
134             return output
135
136
137

```

## A.3 herit.py

```

1
2 def estimate(y, lik, K, M=None, verbose=True):
3     from numpy_sugar.linalg import economic_qs
4     from numpy import pi, var, diag
5     from glimix_core.glmm import GLMMExpFam
6     from glimix_core.lmm import LMM
7     from limix._data._assert import assert_likelihoood
8     from limix._data import normalize_likelihoood,
conform_dataset
9     from limix.qtl._assert import assert_finite
10    from limix._display import session_block, session_line
11    lik = normalize_likelihoood(lik)
12    lik_name = lik[0]
13    with session_block("Heritability analysis", disable=not
verbose):
14        with session_line("Normalising input...", disable=not
verbose):
15            data = conform_dataset(y, M=M, K=K)
16            y = data["y"]
17            M = data["M"]
18            K = data["K"]
19            assert_finite(y, M, K)
20            if K is not None:

```

```
21         # K = K / diag(K).mean()
22         QS = economic_qs(K)
23     else:
24         QS = None
25     if lik_name == "normal":
26         method = LMM(y.values, M.values, QS, restricted=True
27     )
28         method.fit(verbose=verbose)
29     else:
30         method = GLMMEExpFam(y, lik, M.values, QS, n_int=500)
31         method.fit(verbose=verbose, factr=1e6, pgtol=1e-3)
32     g = method.scale * (1 - method.delta)
33     e = method.scale * method.delta
34     if lik_name == "bernoulli":
35         e += pi * pi / 3
36     v = var(method.mean())
37     return g , v , e
```

# Bibliography

- Abadi, Martín et al. (2015). *TensorFlow: Large-Scale Machine Learning on Heterogeneous Systems*. Software available from tensorflow.org. URL: <https://www.tensorflow.org/>.
- Abadi, Martín et al. (2016). “TensorFlow: A System for Large-scale Machine Learning”. In: *Proceedings of the 12th USENIX Conference on Operating Systems Design and Implementation*. OSDI’16. Savannah, GA, USA: USENIX Association, pp. 265–283. ISBN: 978-1-931971-33-1. URL: <http://dl.acm.org/citation.cfm?id=3026877.3026899>.
- Albrecht, Theresa et al. (2011). “Genome-based prediction of testcross values in maize”. In: *Theoretical and Applied Genetics* 123.2, p. 339.
- Alonso-Blanco, Carlos et al. (2016). “1,135 genomes reveal the global pattern of polymorphism in *Arabidopsis thaliana*”. In: *Cell* 166.2, pp. 481–491.
- Angermueller, Christof et al. (2016). “Deep learning for computational biology”. In: *Molecular systems biology* 12.7, p. 878.
- Azodi, Christina B et al. (2019). “Benchmarking algorithms for genomic prediction of complex traits”. In: *bioRxiv*, p. 614479.
- Bernardo, Rex and Jianming Yu (2007). “Prospects for genomewide selection for quantitative traits in maize”. In: *Crop Science* 47.3, pp. 1082–1090.
- Che, Ronglin et al. (2014). “An adaptive permutation approach for genome-wide association study: evaluation and recommendations for use”. In: *BioData mining* 7.1, p. 9.
- Chollet, François et al. (2015). *Keras*. <https://keras.io>.
- Collette, Andrew (2013). *Python and HDF5*. O’Reilly.
- Crossa, José et al. (2010). “Prediction of genetic values of quantitative traits in plant breeding using pedigree and molecular markers”. In: *Genetics*.

- Crossa, José et al. (2017b). "Genomic selection in plant breeding: Methods, models, and perspectives". In: *Trends in plant science*.
- Crossa, José et al. (2017a). "Genomic selection in plant breeding: methods, models, and perspectives". In: *Trends in plant science* 22.11, pp. 961–975.
- Cuevas, Jaime et al. (2019). "Deep Kernel for Genomic and Near Infrared Predictions in Multi-environment Breeding Trials". In: *G3: Genes, Genomes, Genetics* 9.9, pp. 2913–2924.
- De Los Campos, Gustavo et al. (2009). "Predicting quantitative traits with regression models for dense molecular markers and pedigree". In: *Genetics* 182.1, pp. 375–385.
- Desta, Zeratsion Abera and Rodomiro Ortiz (2014). "Genomic selection: genome-wide prediction in plant improvement". In: *Trends in plant science* 19.9, pp. 592–601.
- Dos Santos, Cicero and Maira Gatti (2014). "Deep convolutional neural networks for sentiment analysis of short texts". In: *Proceedings of COLING 2014, the 25th International Conference on Computational Linguistics: Technical Papers*, pp. 69–78.
- Fan, Jianqing, Fang Han, and Han Liu (2014). "Challenges of big data analysis". In: *National science review* 1.2, pp. 293–314.
- Gianola, Daniel (2013). "Priors in whole-genome regression: the Bayesian alphabet returns". In: *Genetics* 194.3, pp. 573–596.
- Gianola, Daniel and Guilherme JM Rosa (2015). "One hundred years of statistical developments in animal breeding". In: *Annu. Rev. Anim. Biosci.* 3.1, pp. 19–56.
- Gianola, Daniel et al. (2009). "Additive genetic variability and the Bayesian alphabet". In: *Genetics* 183.1, pp. 347–363.
- González-Camacho, JM et al. (2012). "Genome-enabled prediction of genetic values using radial basis function neural networks". In: *Theoretical and Applied Genetics* 125.4, pp. 759–771.
- González-Camacho, Juan Manuel et al. (2016). "Genome-enabled prediction using probabilistic neural network classifiers". In: *BMC genomics* 17.1, p. 208.

- González-Camacho, Juan Manuel et al. (2018). "Applications of machine learning methods to genomic selection in breeding wheat for rust resistance". In: *The plant genome* 11.2.
- Grinberg, Nastasiya F, Oghenejokpeme I Orhobor, and Ross D King (2018). "An Evaluation of Machine-learning for Predicting Phenotype: Studies in Yeast, Rice and Wheat". In: *BioRxiv*, p. 105528.
- Habier, David et al. (2011). "Extension of the Bayesian alphabet for genomic selection". In: *BMC bioinformatics* 12.1, p. 186.
- Hayes, Ben and Mike Goddard (2010). "Genome-wide association and genomic selection in animal breeding". In: *Genome* 53.11, pp. 876–883.
- Hayes, BJ, ME Goddard, et al. (2001a). "Prediction of total genetic value using genome-wide dense marker maps". In: *Genetics* 157.4, pp. 1819–1829.
- (2001b). "Prediction of total genetic value using genome-wide dense marker maps". In: *Genetics* 157.4, pp. 1819–1829.
- He, Kaiming et al. (2016). "Deep residual learning for image recognition". In: *Proceedings of the IEEE conference on computer vision and pattern recognition*, pp. 770–778.
- Heffner, Elliot L, Jean-Luc Jannink, and Mark E Sorrells (2011). "Genomic selection accuracy using multifamily prediction models in a wheat breeding program". In: *The Plant Genome* 4.1, pp. 65–75.
- Heffner, Elliot L et al. (2010). "Plant breeding with genomic selection: gain per unit time and cost". In: *Crop science* 50.5, pp. 1681–1690.
- Hirschhorn, Joel N. and Mark J. Daly (2005). "Genome-wide association studies for common diseases and complex traits". In: *Nature Reviews Genetics* 6.2, pp. 95–108. ISSN: 1471-0064. DOI: [10.1038/nrg1521](https://doi.org/10.1038/nrg1521). URL: <https://doi.org/10.1038/nrg1521>.
- Jiang, Yong and Jochen C Reif (2015). "Modeling epistasis in genomic selection". In: *Genetics* 201.2, pp. 759–768.
- Kang, Hyun Min et al. (2010). "Variance component model to account for sample structure in genome-wide association studies". In: *Nature genetics* 42.4, p. 348.

- Korte, Arthur and Ashley Farlow (2013a). "The advantages and limitations of trait analysis with GWAS: a review". In: *Plant methods* 9.1, p. 29.
- (2013b). "The advantages and limitations of trait analysis with GWAS: a review". In: *Plant methods* 9.1, p. 29.
- Korte, Arthur et al. (2012). "A mixed-model approach for genome-wide association studies of correlated traits in structured populations". In: *Nature genetics* 44.9, p. 1066.
- Lan, Kun et al. (2018). "A survey of data mining and deep learning in bioinformatics". In: *Journal of medical systems* 42.8, p. 139.
- Li, Bo et al. (2018). "Genomic prediction of breeding values using a subset of SNPs identified by three machine learning methods". In: *Frontiers in genetics* 9, p. 237.
- Li, Jia-Yang, Jun Wang, and Robert S Zeigler (2014). "The 3,000 rice genomes project: new opportunities and challenges for future rice research". In: *GigaScience* 3.1, p. 8.
- Lippert, Christoph et al. (2014). "LIMIX: genetic analysis of multiple traits". In: *bioRxiv*. DOI: [10.1101/003905](https://doi.org/10.1101/003905). eprint: <https://www.biorxiv.org/content/early/2014/05/22/003905.full.pdf>. URL: <https://www.biorxiv.org/content/early/2014/05/22/003905>.
- Litjens, Geert et al. (2017). "A survey on deep learning in medical image analysis". In: *Medical image analysis* 42, pp. 60–88.
- Ma, Wenlong et al. (2017). "DeepGS: Predicting phenotypes from genotypes using Deep Learning". In: *bioRxiv*, p. 241414.
- Mamoshina, Polina et al. (2016). "Applications of deep learning in biomedicine". In: *Molecular pharmaceutics* 13.5, pp. 1445–1454.
- Martini, Johannes WR et al. (2017). "Genomic prediction with epistasis models: on the marker-coding-dependent performance of the extended GBLUP and properties of the categorical epistasis model (CE)". In: *BMC bioinformatics* 18.1, p. 3.
- Mckinney, Brett and Nicholas Pajewski (2012). "Six degrees of epistasis: statistical network models for GWAS". In: *Frontiers in genetics* 2, p. 109.

- McKinney, Wes (2010). “Data Structures for Statistical Computing in Python”. In: *Proceedings of the 9th Python in Science Conference*. Ed. by Stéfan van der Walt and Jarrod Millman, pp. 51–56.
- Min, Seonwoo, Byunghan Lee, and Sungroh Yoon (2017). “Deep learning in bioinformatics”. In: *Briefings in bioinformatics* 18.5, pp. 851–869.
- Montesinos-López, Osval A et al. (2019a). “A benchmarking between deep learning, support vector machine and Bayesian threshold best linear unbiased prediction for predicting ordinal traits in plant breeding”. In: *G3: Genes, Genomes, Genetics* 9.2, pp. 601–618.
- (2019b). “New deep learning genomic-based prediction model for multiple traits with binary, ordinal, and continuous phenotypes”. In: *G3: Genes, Genomes, Genetics* 9.5, pp. 1545–1556.
- Ogut, Joseph O, Hans-Peter Piepho, and Torben Schulz-Streeck (2011). “A comparison of random forests, boosting and support vector machines for genomic selection”. In: *BMC proceedings*. Vol. 5. 3. BioMed Central, S11.
- Oliphant, Travis E (2006). *A guide to NumPy*. Vol. 1. Trelgol Publishing USA.
- Purcell, Shaun et al. (2007). “PLINK: a tool set for whole-genome association and population-based linkage analyses”. In: *The American journal of human genetics* 81.3, pp. 559–575.
- Qiu, Zhixu et al. (2016). “Application of machine learning-based classification to genomic selection and performance improvement”. In: *International Conference on Intelligent Computing*. Springer, pp. 412–421.
- R Core Team (2019). *R: A Language and Environment for Statistical Computing*. R Foundation for Statistical Computing. Vienna, Austria. URL: <https://www.R-project.org/>.
- Rampasek, Ladislav and Anna Goldenberg (2016). “Tensorflow: Biology’s gateway to deep learning?” In: *Cell systems* 2.1, pp. 12–14.
- Ritchie, Marylyn D and Kristel Van Steen (2018). “The search for gene-gene interactions in genome-wide association studies: challenges in abundance of methods, practical considerations, and biological interpretation”. In: *Annals of translational medicine* 6.8.

- Segura, Vincent et al. (2012). "An efficient multi-locus mixed-model approach for genome-wide association studies in structured populations". In: *Nature genetics* 44.7, p. 825.
- Seren, Ümit et al. (2016). "AraPheno: a public database for Arabidopsis thaliana phenotypes". In: *Nucleic acids research*, gkw986.
- Siva, Nayanah (2008). *1000 Genomes project*.
- Storey, John D. and Robert Tibshirani (2003). "Statistical significance for genomewide studies". In: *Proceedings of the National Academy of Sciences* 100.16, pp. 9440–9445. ISSN: 0027-8424. DOI: [10.1073/pnas.1530509100](https://doi.org/10.1073/pnas.1530509100). eprint: <https://www.pnas.org/content/100/16/9440.full.pdf>. URL: <https://www.pnas.org/content/100/16/9440>.
- Togninalli, Matteo et al. (2017). "The AraGWAS Catalog: a curated and standardized Arabidopsis thaliana GWAS catalog". In: *Nucleic acids research* 46.D1, pp. D1150–D1156.
- Van Rossum, Guido and Fred L Drake Jr (1995). *Python tutorial*. Centrum voor Wiskunde en Informatica Amsterdam, The Netherlands.
- Webb, Sarah (2018). "Deep learning for biology". In: *Nature* 554.7693.
- Zhang, Zhiwu et al. (Mar. 2010). "Mixed linear model approach adapted for genome-wide association studies". In: *Nature Genetics* 42.4, pp. 355–360. DOI: [10.1038/ng.546](https://doi.org/10.1038/ng.546). URL: <https://doi.org/10.1038/ng.546>.
- Zhou, Xiang and Matthew Stephens (June 2012). "Genome-wide efficient mixed-model analysis for association studies". In: *Nature Genetics* 44.7, pp. 821–824. DOI: [10.1038/ng.2310](https://doi.org/10.1038/ng.2310). URL: <https://doi.org/10.1038/ng.2310>.